

Automated negotiation in the game of diplomacy - Report 3: Overview and Software Validation

Matthias Hueser, Andras Slemmer, Luca Deltodesco, Luke Tomlin, Cliff Sun

Monday 12th December, 2011

1 Project synopsis

The Game of Diplomacy is a round-based multiplayer strategy-game set in post WWI Europe. Each player controls a power which initially operates from a number of home-bases. The ultimate objective of each player is to control the majority of supply-centers on the map. A game is structured into a number of rounds, in which each player formulates a move order to conquer additional supply centres or defend existing ones. Besides the necessity to create effective long-term strategies negotiation is integral to the game-experience. There is a thriving internet-based community around the Game of Diplomacy and many AI players, so called bots, have been created in the past. Most of the existing AI players are quite effective decision-makers in general but few of them support negotiation or reasoning about the relationship of the powers in the game. Our project aims to create a framework for Diplomacy, including a GUI client, a game server and a collection of bots. Crucially the highest evolution of our series of bots should be able to analyze and act upon messages of human / AI players and issue messages to other players. Finally experiments are conducted to determine whether there is a reward associated with negotiation capabilities. This will require careful study and implementation of automated negotiation techniques from other domains in AI/Game Theory. Thus the project gives us ample opportunity to be creative and test different approaches using the DAIDE-conforming game framework also produced for the project. DAIDE is a standard client-server protocol which serves as a 'common language' for different Diplomacy bot implementations, enabling them to play a game together on a conforming server.

1.1 Requirements engineering / management

High-level requirements were discussed in the initial phase of the project (weeks 1/2) with our project supervisor – Iain Philipps – who is our customer in Software Engineering terms. The terms of the collaboration were such that only certain high-level goals were defined and decisions about which AI techniques to use were left with the team-members.

- The over-arching goal and hence the most critical requirement is to design and implement a negotiation bot for Diplomacy whose performance is then compared to existing bots in an experimental setting.
- A collection of other bots which are not capable of communication should be created for purposes of comparison. These should use proven techniques from the fields of Game Theory and Multi-Agent systems, including learning, tactics and long-term strategies. All bots should support the DAIDE protocol to compete against existing bots on the DAIDE windows server.
- To facilitate experimentation an open-source framework conforming to the existing DAIDE protocol should be created. This framework comprises a server which hosts games for automated and human clients and a tool to gather game statistics. The imperative behind this is to feed-back into the Diplomacy community as a whole, providing a multi-platform server (Windows DAIDE server exists) for the first time.
- In addition a framework for simple generation of Diplomacy bots is envisioned: Right now a player needs to re-invent the wheel by implementing well-established game-tree search techniques or certain Machine Learning paradigms in low-level terms. This is clearly unsatisfactory. Our framework helps to shift the focus to interesting new techniques from academia which can be coded up and layered on top of primitives. Actually this framework is used for the 2nd objective above to quickly generate a series of bots with increasing functionality. (Abstract client pattern) This also works as a proof-of-concept for the server framework. At the users convenience extensions are written in the high-level language Haskell which is more expressive than C++ and natural to practitioners in the field of Artificial Intelligence / Game Theory.
- Finally to test our server and provide a visually pleasing user interface we will deliver a DAIDE-conforming GUI client for the Game of Diplomacy. Rather than forcing the player to use a particular device all back-ends of the HaXe platform are supported, including HTML5/JS, Flash and C++.

In the above listing the requirements are decreasing in importance. Some of the above look indivisible in nature but we have still managed to extract user-stories to support time-boxed iterations. Requirements management in Weeks 3-11 was two-fold: Each team-member kept in mind the overall contract with our supervisor (see above) which defines the direction of the project. Based on this User-Stories were defined which guide the primary development focus during the iterations.

1.2 Change / Risk management

We recognized that change is inevitable in such an ambitious and multi-faceted project. Therefore we took great care in defining our contracts carefully to split up work. Actually this was

helped substantially by the existence of the DAIDE protocol. Through its rigorous specification it forced everybody to program to an existing interface and there was little scope for confusion as to what needs to be supported.

Each of the high-level objectives outlined above can be achieved separately and as a result there are little dependencies between team members. If communication between two components fails it is matter of determining which are not properly realizing the DAIDE protocol. A welcome side-effect is that we avoid endless debugging sessions to make two parts of the project work together. They just need to be tested against a DAIDE reference.

Using the 'AI generation' framework explained above implicitly makes coding of AI techniques incremental. For instance if a well-tested game-tree search or Learning algorithm is in place, each evolutionary bot can delegate these tasks without worrying about how they are implemented. Turning this into a reality required defining a layered design in Weeks 3-5 which all future work respected.

1.3 Planning / task estimation

All planning took place in the weekly group meetings (for the structure of a meeting see example at the bottom of document). A general rule for new user-stories was that they relate to the global objectives defined in the contract with our supervisor (see above). This avoided falling ill with 'featuritis' at any point in the project.

In the AI realm a new user story was typically introduced by a team member. Generally we shied away from implementing exotic algorithms that are not clearly understood by all team members, adopting the "Kiss principle". Also there should be a balance between different paradigms in the AI design, which enables us to make meaningful experiments for the final report. For instance it is not seen as fruitful to performance-tune beyond the point where much value is generated on the whole. Rather the focus should be on negotiation. Associated with each user story we started a quick poker game to estimate the required time-resources for completing the feature. Obviously members with more experience in a domain were given more weight in the decision process. We usually arrived at a decision about a future iteration unanimously.

For the server the process was less creative and controlled by the need to conform to the DAIDE protocol. Conversely planning was easier because the process of creating a parser is understood by all team members.

1.4 Progress metrics

There are multiple ways in which we can measure the progress of our AI. We adopted an aspect-oriented approach in measuring, recognizing that some requirements are fulfilled qualitatively and others can be stated in terms of numbers.

- A natural but informal progress metric is to simply estimate for each high-level objective

defined in the contract a percentage of the features which have been successfully implemented. This places emphasis on the big picture rather than measuring the performance of a particular bot we have created.

- Another metric which only pertains to the AI side-of-things is a quantitative measure of how do we fare against typical existing bots. Such a metric could be gathered for example in a Round-Robin tournament. As a canonical example we adopted the existing DumbBot which despite its name has considerable playing strength. According to our supervisor systematically out-performing DumbBot presents the technical benchmark for the project.
- Relevant to the server-framework and the GUI client we can define a measure by the percentage of valid DAIDE messages supported. Once this is close to 100 % it immediately follows that both are DAIDE-conforming, one of our objectives outlined above.
- As a GUI is hard to test automatically we left 10 mins of our weekly discussions for informal game walkthroughs. All team members judged how natural / visually pleasing the interface was and what features could be supported in the next generation. We have not used any formal methods here but trusted our experience with playing similar strategy games.

1.5 Detailed AI metrics

Within each of these milestones, we can weigh a bot's ability using its only application - playing the game of diplomacy. For this, we need to come up with a set of metrics with which we can evaluate a player's performance in a game. We are planning to equip the server with a tool to measure these statistics during the game. By combining these indicators in a weighted fashion (possibly with some needed calibration), we should be able to compute an accurate score of how well a particular AI is playing. This can be validated by measuring the correlation with a high-score and the number of games won/lost.

Some key indicators are, with some having precedence over others:

1.5.1 Games won/lost

This is clearly the most important metric and overrides all others. However it gives little insight into what actually caused a bot to lose or win the game.

1.5.2 Supply centres controlled

With regards to supply depots, the winning player will own half of them (eighteen), with the second-most successful player owning the second highest amount, and so on and so forth. A player with no supply-depots is very close to becoming a losing player.

1.5.3 Units lost during the game

Units lost is a difficult metric to quantise - whilst it may immediately seem that losing many units is a bad thing, these could be due to tactical masteries involving trapping and deluding many foes. Objectively, it is probably advisable to refrain from losing units where possible.

1.5.4 Provinces conceded during the game

Similarly, conceding provinces, unless done in a planned fashion, is a general indicator that a player is not performing well.

1.5.5 Negotiation 'strength'

If it were possible to evaluate other players' "attitudes" towards the AI player, one might be able to deduce the competence of a bots negotiating skills. For instance, making enemies is widely regarded as a poor move, especially if those enemies are actively hostile against you. A better tactic might be to give them the illusion of friendship, whilst aligning them for a backstabbing manouever. If we could create a way to reliably score the subtleties of negotiation between players, it might aid us in the creation of more advanced diplomising AIs.

1.6 Team velocity / Milestones

We can define our team velocity as the positive change in progress metrics overall. This means that we can allocate team members to the metric which currently has priority. These are typically the ones that are covered by the user-stories for the current iteration. Having said that we strived to progress in each part of the project to discover problems and dependencies we have not anticipated early.

Whereas the velocity-scheme served as a tool to measure our progress internally we communicated our progress to the supervisor in terms of coarse-grained milestones. This had the advantage that we did not overload our customer with technical details that only we as developers care about. Also they co-incided with the frequency of meetings with our supervisor. The milestones agreed upon were to create the AI framework and the server, and then proceed to iterate a build of an AI, progressively increasing its abilities and making it better.

1.7 Description of iterations per subsystem

1.7.1 Framework and Server

For the server we can define some functionality 'chunks' which are strictly independent from others, these guided how we allocated user stories to the 3 iterations outlined below:

- First and foremost the server needs to conform to the DAIDE specification which is an issue of setting up a connection between server and client and parsing conforming issues.

This is a well understood task from the 2nd years Compilers course and hence it formed our first iteration. At the end of the iteration substantial testing was performed (see section on Testing for details). Notice there is no dependancy on any other component being working.

- Secondly the server needs to advance the game state in response to valid DAIDE messages. This was a longer iteration since all special cases treated in the game rules needed to be catered for. Again existing bots and GUI clients could be used to instrument the server and validate correct reaction.
- The last iteration of the server consisted of User-Stories which were created during the weekly meetings. They were not strictly necessary but but useful utility features that served to round up the product. Free time was now spent on refactoring the code and fixing remaining bugs.

1.7.2 AI framework / bots

These iterations have run concurrently with server developments (3 iterations detailed above):

The user stories we have proposed largely correspond to ideas DAIDE wiki and research into machine learning / search techniques.

Each user story contains an iteratively improved version of our current benchmark bot. Hence the overall AI gains functionality and depth on each iteration. While there are some components which every bot should have (for instance a forward-search in the game state-space) specifics are discussed during the group meetings. The idea is that each team member involved in the AI has done some research and proposes new AI techniques which shall be implemented in the next time-box.

Then once a feature is judged mature we incorporate it into the generic AI framework discussed at the beginning. This extends the code-base we treat as given for the next iteration.

- The first iteration – HoldBot – is simply an AI player that does not impede the game progress. It always performs the same move – it holds. It can be used as a simple unit test for the server framework.
- As a first refinement is RandomBot is created, which now considers every move defined in the Diplomacy rules and selects one uniformly at random. Reflection capability is optional, that is the bot may or not reason about how the random moves affected its standing in the game.
- The first bot designed to compete with other bots through the DAIDE protocol is StrategyBot. At a high-level it must use techniques for performing a search in the game tree, formulate general strategies. Also it should be able to improve its performance by analyzing previous strategies or moves.

- The last iteration is the NegotiatingBot which additionally exchanges messages with other players. It needs to have an internal model of the other players intentions and reason about which is friend or foe.

1.8 Progress report for Week 10/11

Unfortunately we suffered some setbacks in the beginning with planning and evaluation of how long it would actually take us to complete the server and framework. These miscalculations have now been addressed and we are progressing with the AI creation while improving server / framework as appropriate.

On the AI-side HoldBot and RandomBot are released (to the current specification of the AI framework). The current iteration focusses on implementing the features of StrategyBot.

2 General validation

2.1 Server testing

- A very practical and simple method for us to test the server is to let RandomBots join a game and issue arbitrary move commands. We estimate that we get good coverage using this method since a random bot is expected to issue every possible move over a large number of games. We expect that the server does not crash and advances the game state correctly. We check successful completion by checking for any error return codes / exceptions thrown by the server. The latter is tested by cross-checking the game state with the Windows DAIDE server which we forward all messages that our server receives. Since both observe the same game the messages send to clients should not differ. To automate the tests, we created a script that intercepts client-server communication and then invokes 'diff' on the messages.
- To inspire further confidence in the code's correctness we use the Quick Check tool developed for the Haskell platform: Given pure Haskell code (that is containing no I/O) it generates random function inputs and theoretically allows us to test a function exhaustively. For instance to test a parsing function we can specify the invariant that it accepts precisely the members of the DAIDE language and returns a parsing error otherwise. For the 'state advancing logic' of the server we can check certain properties that we know to be true in general. An example would be: A supply centre can only be conquered if an army of the player was sent there. We decided against a parallel module hierarchy for the tests but instead defined submodules of the form 'X_quickcheck'. This allowed us to use the tests as documentation when writing the functions and avoided synchronization issues. This also facilitates testing of functions that are not exported from a module.

- To augment testing with QuickCheck one team member has encoded a number of representative game situations as unit test cases. Largely these were taken from the DiplomacyRuleBook, one of the references that specifies correct state advancement. The unit tests were written in HUnit, whose feature-set largely parallels JUnit for Java programming. The whole testing framework can be automated, tests being packages hierarchically in test suite, modules and collections. As QuickCheck tests they live in a module sub-ordinate to the functions tested.

2.2 AI client testing

- A very simple test is to determine whether an AI client does not impede the progress of the game. This can simply be tested by letting the client on either the functional Haskell or the existing DAIDE server (Windows). The script running the test will intercept any error messages issued to the client and deem the test to be failed if any syntax errors (no valid DAIDE message) or semantic errors (moves that do not make sense with respect to the current game-state) were flagged.
- Anything that goes beyond liveness and absence of errors cannot be tested trivially. Instead of testing we let bots play in Round-Robin tournaments, collect game statistics and assign each bot a playing strength metric. The exact way and the ingredients of the formula were outlined in the section on 'AI metrics'. Ultimately playing strength equates to faring robustly against a large collection of existing bots.
- Some white-box testing was done to determine how the AI arrives at a decision and check if its reasoning is sound. We designed the different parts of our AI with testability in mind, often using very tiny functions. As an example consider an interface that measures the utility of a game state to a certain player. It can be tested by simply comparing the result with our pen-and-paper calculations. Some of these tests were coded up as HUnit test suites.

2.3 GUI client testing

- For the GUI client similar criteria to the AI client applied. The client should be live during the game and not issue mal-formed messages. A simple script that instruments the GUIs command-line interface was put in place for this.
- Unit testing in the HaXe proceeded as recommended in its manual: Simple test cases to check correct parsing of SVG files (encode the map topology) were created by subclassing a TestCase and putting appropriate checks in place.
- Functional requirements were largely in-tangible, such as presenting a intuitive, visually appealing interface to the user. Our division of labour allowed that team members involved in the AI could give feedback. Roughly 15 mins of each meeting were dedicated

to discuss the user interface: flagging new issues and tracking progress on fixing previously discovered bugs. These tests were conducted for each end-user device targeted by the HaXe framework, such as mobile phones or internet browsers.

- While there exist tools for automated-checking of web interfaces the team dismissed this as misguided effort better spent elsewhere. The marginal utility would be to know that all buttons worked correctly which can also be tested by manual inspection.

2.4 Examples of bugs discovered during testing

- Grave errors when parsing certain DAIDE messages. This had to do with the ambiguity of the DAIDE grammar and a work-around needed to be found (discovered through QuickCheck testing)
- Inconsistences on the server-side. This proved a pain to rectify and to avoid such problems in the future we decided to come up with a robust scheme for concurrency in the server code. (discovered through sporadic server stress testing)

2.5 Static code checking / tools

All deliverables, including code and documentation was produced using Emacs / Gedit or other similar simple text editors. We have not used an IDE but several Emacs modes and extensions (directory browsing) helped us to keep track of the overall structure of the codebase. Our strategy to achieve good coding style was readability was two-fold: Rather than making up our own coding standard we have adopted the programming guidelines from “<http://www.haskell.org/haskellwiki/Programming.guidelines>”. Secondly each team member was required to run each (compiling) commit through ‘HLint’ and the ‘Haskell Style Scanner’ and resolve issues if necessary.

2.6 Code documentation

We treated documentation as a deliverable only second in priority to working code. The coding standards defined at the beginning of the project detailed the style of documentation and good practices to avoid running-out-of-sync issues.

2.6.1 Developer documentation

Developer documentation consisted of JavaDoc-style comments for each data type, typeclass (interface) and function documentation. The particular tool used is “Haddock” which allows generation of API reference pages (HTML).

2.6.2 User documentation

Besides documenting exported API through “Haddock” (primarily the AI framework for creating Diplomacy Bots) we agreed that tutorials need to be written to explain typical use-cases of the library. Only then can we hope that our library is re-used by other members of the Diplomacy community which was an objective defined at the outset of the project. Team members agreed that once we go public with our project (BSD open-source licensed) we also need to have a Sourceforge website.

2.7 Code inspection / Refactoring

Each function of the code is ideally represented in three different ways: as executable code, unit test and exported “Haddock” comments. For some trivial cases one of the latter might be missing. This approach has simplified code inspection and walkthroughs tremendously. In a typical ‘design-check’ session which is conducted roughly fort-nightly a team-member will check a particular code in detail. This involves clarifying comments, naming of functions and checking overall soundness of design. Often major refactorings were suggested as it was discovered that pieces of code were conceptually similar and hence could be extracted into a common module. Not only were we challenging our tendency to ‘repeat ourselves’ in the code but also could spread knowledge about the code-base in the entire team. By the regularity of refactoring / design effort we wanted to avoid that the code-base grows out-of-shape and acquires “technical debt”.

2.8 Acceptance tests

Acceptance testing has not yet been finalized. In a recent meeting with our supervisor we conducted a short walkthrough of the GUI to gather feedback. The functionality we have implemented was received well, with some minor changes suggested. Final integration / system testing will be done during the holidays and the first week of the Spring term. We aim to keep our supervisor involved by sending regular screenshots of the GUI and results from the Bot-tournaments.

2.9 Stress testing

2.9.1 Server loading

Our server can be stressed by increasing its load, flooding it with malformed messages in spirit of Denial-of-Service attacks. The overall aim is to find a threshold load under which our server just crashes. Another test is to connect multiple AI clients that issue conforming DAIDE messages at a rapid pace. The question then is whether the server still advances the game state correctly. If this is not the case we would need to check the server code for possible race conditions explicitly. Initial measurements indicate sufficient robustness.

2.9.2 AI / GUI client loading

In a similar fashion GUI and AI client can be run against servers issuing malicious messages at a rapid pace. We allow arbitrary behaviour in this case, but the client should still discover the problem and disconnect safely. Similarly we have to defend against an adversary who addresses spam to force our player to time-out during move generation.

3 Managerial Documentation

3.1 Collaboration tools

We discovered that the best tool was actually gathering in front of a white-board (we did not yet have IdeaPaint TM), sticking some user story cards and talking through the overall design. Simplified UML diagrams and examples were used to elaborate design alternatives at our disposal. Anything that was of permanent value to the team was usually recorded in a logbook (see example of a LogBook page below). For this purpose we assigned a member to take minutes of each meeting and create a digest that would be put into the issue-tracking system of GitHub. Other resources of GitHub, such as progress, issue flagging and milestone setting were used sporadically. Next time we would probably use a more sophisticated tool for resource management, enabling us to categorize research papers, coding standards, progress reports in a better way.

3.2 Terms of collaboration

3.2.1 Working hours / patterns

As suggested in the spec each group member spent the rough equivalent of one working day per week on the project. We assured there is a good mix of work to be done by each member, so that nobody will become bored or disinterested by tedious work. Of course some long coding sessions were needed to fix bugs and integrate components. We tried to work against such problems though by having regular meetings where the design and any bugs are discussed.

3.2.2 Group meeting conditions

Each group meeting was held in the conference room opposite the labs. We ensured that each member was free for the day afterwards to cater for the fact that meetings sometimes had to overrun the expected time. Whiteboards were used to discuss the design with UML / package diagrams.

3.2.3 Group meeting structure

Besides discussing user-stories and mini-milestones for the next iteration group meetings were used to do the following things

- Assessing risks early: For instance we dedicated one half-session to come up with a technique to cope with race conditions occurring on shared data structures in the server. Two team members had alternative solutions to this solution with the group finally agreeing with using Software Transactional Memory, which is natively supported in the Haskell GHC extensions. Another technical risk was using the DAIDE system, which is legacy in nature (ambiguities in the grammar of DAIDE messages etc.). The team had to discuss how to proceed in these cases.
- Defining testing strategies: As shown in the testing section we have used different techniques in some are more appropriate in certain cases. Usually the team member working on the code made a suggestion and argued their case. We reckoned that everybody had a stake in the decision because the code was usually tested / reviewed by a person different from the coder. Issues of code coverage and general confidence into the code / testing suite were also raised here. The overall aim was to make out weak spots that need urgent clean-up and improvement.
- At the beginning of each iteration, we plan out the work distribution that needs to be done. If more than one person is writing code we need to raise awareness of dependencies that we may have between these members. These could take the form of overlapping code or module dependencies. Before any coding starts the respective members need to agree on a common strategy to avoid code duplication and factoring common parts out.
- Each second meeting we try to do some “meta-learning” - that is what went wrong and how can we improve our process in the following week. For instance it was discovered early that different knowledge bases about Haskell / monads could prove an obstacle to collaboration. Hence we agreed on putting relevant research papers on GitHub to educate all team members in these central concepts. These discussion usually provide the basis for what appears in the reports on “Software Engineering”.

3.2.4 Record of group meetings

Meeting Type	Date	Duration	Iteration	Summary of Meeting
Supervisor	12/10/2011	1 hr	1	Initial meeting with supervisor, discussing the scope of the project and confirming technologies that we were going to use.
Team	14/10/2011	2 hr	1	Iteration 1 first meeting and first meeting with group, discussing division of work and background reading. In addition we finalised our decision on what technologies to use and started to test compatibilities between them (Haskell and Haxe). Also aim is to learn these new technologies from other members/resources on the internet.
Team	21/10/2011	2 hr	1	Mid-Iteration Meeting to see how the server and the client were coming along. Server was to adopt the Daide protocol which provides a framework for writing a Diplomacy AI Bot and support for allowing other Daide AI Bot's to play on the server.
Team	28/10/2011	1 hr	1	End-of-Iteration 1 meeting. Client and Server can send messages to each other
Team	28/10/2011	1 hr	2	Beginning-of-Iteration 2 meeting. Goals for end of iteration are to adding parsing for client and server side to parse press level 0 daide messages (basic diplomacy game mechanics). Also try and get the order resolution done (game state changes).
Supervisor	4/11/2011	1 hr	2	Second meeting with supervisor, discussing any problems we had and how we planned to continue with building the server and the AI. In addition discussed how we would go about implementing negotiation into the game.
Team	4/11/2011	1 hr	2	Mid-Iteration Meeting for Iteration 2 to discuss problems, mainly that parsing on the server is taking longer than expected due to amount of coursework and interviews group members have had
Team	11/11/2011	1 hr	2	End-of-Iteration 2 meeting. Goals not met this week due to the high amounts of other work that group members have (see previous). Parsing still not complete and order resolution not implemented still.
Team	11/11/2011	1 hr	3	Beginning-of-Iteration 3 meeting. Aims for this iteration adjusted to account for targets that were missed in previous iteration. Parsing to be completed create a framework which handles communication for AI bots, and provides a base layer for AI bots. Initially create a basic AI bot that just holds (holdbot). GUI is to be improved with a game map defined by clickable areas.
Supervisor	18/11/2011	1 hr	3	Third meeting with supervisor, discussing the development of the AI and whether targets for the project can be met.
Team	18/11/2011	1 hr	3	Mid-Iteration Meeting for Iteration 3. Bot framework and Holdbot (skelbot) under development. Map for GUI also making good progress.
Team	25/11/2011	1 hr	3	End-of-Iteration 3 meeting. Skelbot (and its related components) and Holdbot completed. However can't play on server as the order resolution still not implemented.
Team	25/11/2011	1 hr	4	Beginning-of-Iteration 4 meeting. Aim to create RandomBot which searches game state for all available moves and chooses a random move. Data definitions are to be
Supervisor	02/12/2011	1 hr	4	Discussing plans for working over winter holidays to reach targets and discussing details of 3rd and final reports.
Team	02/12/2011	1 hr	4	Mid-Iteration Meeting for Iteration 4. Randombot
Supervisor	09/12/2011	1 hr	4	Demo of GUI working on windows Daide servers.
Team	09/12/2011	1 hr	4	End-of-Iteration 4 meeting

3.2.5 Log-Book

Our log-book records details progress for the current iteration and any problems / dependencies that need to be resolved. It is primarily standard text but can contain UML / package diagram visualizing central ideas of our design.

Our logbook structure looks like something as follows:

Individual Progress: Here we have a section for each member and we discuss the progress of the work for the current iteration. If it's a mid-iteration meeting then we see how the work for each member is coming along and whether they are able to reach the goals for the iteration. If not (member may have coursework and interviews that week for example), then we try and delegate some of the work to other members who may finish their iteration early. If no one is able to take on extra work from someone else, then we decide which of the goals is more important to achieve for the iteration (which may affect the goals of the next iteration) and place high priority on that.

Andras Progress:

Cliff Progress:

Luke Progress:

Matthias Progress:

Luca Progress:

Overall Progress: If it's a mid-iteration meeting then we see how group members have got on and try and judge if we need to adjust work rates and work distribution to achieve our goals for the iteration. If we find that we may not be able to reach the goals for our current iteration then we put the current state of the iteration as [AT RISK].

If it is a end-of-iteration meeting (and first meeting of the next iteration) then this section serves as a checklist to see which goals we have achieved and to assess whether we are on track to complete the project in time. If any goals have not been achieved then we put these as high priority to complete in our next iteration.

Next Iteration Goals: Only fill this out if it's a beginning-of-iteration meeting we are having. Using the information from above and taking into account the final goals for the project, we come up with the goals for the new iteration. This means coming up with a high-level goal and then breaking it down into work distribution for each individual member. E.g:

Overall: Server and client is able to successfully parse incoming messages from each other at press level 0

Andras: Complete parsing for server at press level 0

Cliff: Help with completing the parsing

Luke: Complete dynamic map representation

Matthias: Help with completing map representation

Luca: Complete parsing for client side

3.3 Version control

The distributed source code control system Git was used for all deliverables - that is unit tests / documentation, tutorials and presentations. We have made it strict policy to avoid both conflicting / non-compiling commits. The former required good coordination between team members about which files should be edited in what time-window. Any file that is machine-

generated should be ignored by the file tracker to avoid confusion and wasted memory space in the repository. For simplicity we avoided working on different development branches. All code was backup-ed in the project directory of DoC at regular intervals.

3.4 Automated build

The CABAL package management system was used as recommended by most Haskell coding standards. The different parts of the project (Server, Client, AI) were packaged separately to avoid annoying compile dependencies and allow team members to work separately on different components. In our judgement a suitable CI server does not exist for Haskell and using one would have been overkill.

3.5 Management / Organisational policies

We tried to keep it lean here: There is a culture of trust within the group which can be justified because we have been rigorous about commit / coding rules at the outset. Sub-optimal design decisions were usually caught in the “design-check” sessions explained above. For some parts of the codebase exotic techniques like Software Transactional Memory were used. This proved problematic because team members usually did not have the time to learn these new concepts, making code review difficult. In such cases ‘research papers’ presenting the necessary techniques were usually sent to each team member involved in the code inspection. We had to take a pragmatic approach here – after all team members sometimes have very different backgrounds in programming paradigms / coding techniques. It is a necessary side-effect of division of labour that not everybody is an expert in everything.

3.6 Knowledge transfer within the group

As previously mentioned, there is a significant amount of knowledge variation within members of the group. Some of us are skilled with Haskell coding, whilst some of us are merely acquainted with it. Similarly, GUI design, AI design, general coding skills and practices all varied. We have remedied this somewhat by working in pairs or more, allowing the parts of the knowledge to diffuse (in a manner similar to osmosis :)) between group members as problems are encountered and solved. Additionally, papers and other learning tools have been shared (via GitHub and Gmail) to aid the code writing process. Where learning-by-application and reading have failed, group members have also been keen to help each other, spending time drawing UML diagrams and explaining concepts in detail. The focal points of these discussions have usually been the weekly group meetings. Logbook allowed us to document what we have learned and avoid introducing bad design / bugs more than once. Last but not least we decided to have a “glossary” section in the log-book where we explain give names to the most important concepts of our design. For instance we have defined the AI framework in terms of metaphors from the Neuroscience domain (Brain etc.). For this to be effective and improve communication every team member needed to be aware of the terms.